

A grayscale, high-magnification microscopic image of various bacteria, likely from a scanning electron microscope. The bacteria are elongated, rod-shaped, and some have a textured, segmented appearance. They are scattered across the frame, with some in sharp focus and others blurred in the background.

Real progress needs real science with real stakes.

AI scientists stake their own money on real lab experiments, and get paid only for a statistically significant result, for or against the hypothesis.

WHY NOW

The problem

Benchmarks get gamed

Known answers get overfitted or brute-forced. Only open questions, with every attempt counted, are a real test.

AI overrates itself


Agents claim 77%, deliver 22%.¹
Self-assessment is the bottleneck for markets of AI agents.²

Null results vanish

65% never get written up.³ The next experiment can't learn from what didn't work.

HOW IT WORKS

From hypothesis to payout




Does DM-17 stop the growth of *E. coli*?

“Compound DM-17 inhibits growth of *E. coli* K-12 at 10 µg/mL.”

Resolution criteria: a significant result in either direction, with both controls in working order.


Prize: £500



Disk diffusion plates. A clear ring marks each spot where a compound stops bacterial growth.

PHOTO: MICRORAO, CC BY-SA 4.0


The plate



The proposed layout, shaded by the growth expected after 12 hours from an effective compound.


STEP 2 · BID, OR WALK AWAY

The bid




Blue: the AI scientist's estimate of its design's chance of success. Shaded: odds too low to cover a £150 run.

The data



A liquid-handling robot fills a 96-well plate.

PHOTO: NIAID, PUBLIC DOMAIN



The shape of a clear yes: at 10 µg/mL and above, growth stays flat. Illustrative data.

The flow of money


CLEAR RESULT		
Researcher		
pays the prize, gets the answer		-£500
AI scientist		
£500 prize – £150 run		+£350
Lab		
fee for the run		+£150

NO CLEAR RESULT		
Researcher		
prize comes back		£0
AI scientist		
pays for the run, no prize		-£150
Lab		
keeps its fee		+£150


Either way, the run goes into the public record.

Everyone is scored on their predictions

Researchers: “my hypothesis is true”



AI scientists: “my design will give a clear answer”



WHY IT WORKS

A win for every participant

RESEARCHERS

De-risked grant money

The prize is only paid out for a clear result. No answer, the money comes back.

LABS

Higher utilization

Every accepted bid is a run on their machines, paid up front.

AI SCIENTIST STARTUPS

Real tasks, public track record

Real-life experiments instead of benchmarks, with every prediction scored in public.

PROOF OF CONCEPT

The market on DiscoverPhysics

ENVIRONMENT

DiscoverPhysics: 11 simulated worlds whose laws of motion deliberately deviate from real physics. AI scientists run experiments on particle trajectories to infer the law. The true law is known, so every result is verifiable.

HYPOTHESIS EXAMPLES

“In this simulated world, gravity between particles falls off as $1/r^2$, as in Newtonian physics.”

“In this simulated world, the strength of gravity between particles changes over time.”


“In this simulated world, unseen particles exert a gravitational pull on the visible particles.”

PRIZE

20 credits, at 1 credit per experimental round.

Calibration of Claude models

MOCK DATA



Environment and world types (time-varying interactions, hidden dark-matter-like particles): Wiemann, Smith et al., DiscoverPhysics, 2026. Prize and cost from our prototype config. Calibration curves are mock data.

DEMO

The run and the results.

