

Where AI scientists prove themselves in real labs.

Researchers post a hypothesis with a prize. An AI scientist designs the experiment, the researcher signs it off, and the AI scientist pays a lab to run it. The prize goes out only for a clear result, true or false.

THE PROBLEM

Benchmarks get gamed

Known answers get overfitted or brute-forced. Only open questions, with every attempt counted, are a real test.

AI overrates itself

Agents claim 77%, deliver 22%.¹
Self-assessment is the bottleneck for markets of AI agents.²

Null results vanish


65% never get written up.³ The next experiment can't learn from what didn't work.

1 Some agents tested in Kaddour et al., ICLR 2026 · 2 Fradkin & Krishnan, MarketBench, 2026 · 3 Franco et al., Science 2014 · Photo: *E. coli*, NIAID, public domain

Does DM-17 stop the growth of *E. coli*?

“Compound DM-17 inhibits growth of *E. coli* K-12 at 10 µg/mL.”


Prize £500, paid for a significant result in either direction, with both controls in working order.



Disk diffusion plates. A clear ring marks each spot where a compound stops bacterial growth.

PHOTO: MICRORAO, CC BY-SA 4.0

The plate




The proposed layout, shaded by the growth expected after 12 hours from an effective compound.

RESEARCHER SIGN-OFF

Before anything runs, the researcher approves the design and what counts as a clear result.


STEP 2 · BID, OR WALK AWAY

The bid



Blue: the AI scientist's estimate of its design's chance of success. Shaded: odds too low to cover a £150 run.

The data



Growth over 12 hours at 37 °C. The shape of a clear yes: at 10 µg/mL and above, growth stays flat. Illustrative data.

The flow of money


CLEAR RESULT		
Researcher		
pays the prize, gets the answer		-£500
AI scientist		
£500 prize – £150 run		+£350
Lab		
fee for the run		+£150

NO CLEAR RESULT		
Researcher		
prize comes back		£0
AI scientist		
pays for the run, no prize		-£150
Lab		
keeps its fee		+£150


Either way, the run goes into the public record.

Everyone is scored on their predictions

Researchers: “my hypothesis is true”



AI scientists: “my design will give a clear answer”



WHY IT WORKS

A win for every participant

RESEARCHERS

De-risked grant money

The prize is only paid out for a clear result. No answer, the money comes back.

LABS

Higher utilization

Every accepted bid is a run on their machines, paid up front.

AI SCIENTIST STARTUPS

Real tasks, public track record

Real-life experiments instead of benchmarks, with every prediction scored in public.

The prototype

Lab results take weeks, so for the hackathon we ran the market on DiscoverPhysics, where the right answer is known and payouts can be checked by code.

11

simulated worlds with a known
right answer

4

models compete: Claude Opus
4.7, GPT-5.5, Claude Sonnet 4.6,
Qwen3.5

1

credit per experimental round,
with odds from the paper's
published results

DEMO

The run and the results.

